



Fintech
Open Source
Foundation

An introduction to the FINOS AI Governance Framework

What is FINOS?



Fintech
Open Source
Foundation

The Fintech Open Source Foundation (FINOS) is an **independent, non-profit membership organization** whose purpose is to **accelerate** collaboration and **innovation in financial services** through Open Source software & standards.

FINOS provides neutral hosting for **global financial system constituents to collaborate in the open**, hosts the **Open Source in Finance Forum**, offers **training and certification** for its project and practices and produces the annual **State of Open Source in Financial Services** report.

Oh, and in case you're wondering, it's pronounced **fee·nowz**



Platinum



100
2025
MEMBERS

Gold



Silver



Associate



We are now truly cross-industry!

See all our
members:



BUY SIDE

WELLINGTON
MANAGEMENT®



BlackRock.



MARKET DATA AND FINANCIAL MARKET INFRASTRUCTURE

S&P Global



FACTSET



MOODY'S

JAPAN SECURITIES
CLEARING CORPORATION

CONSULTING FIRMS



luxoft
A DXC Technology Company



Synechron

BCG BOSTON
CONSULTING
GROUP



publicis
sapient

INDUSTRY ASSOCIATIONS

ISDA
Safe,
Efficient
Markets

ISLA
INTERNATIONAL
SECURITIES LENDING
ASSOCIATION

FIX TRADING
COMMUNITY™
INDUSTRY-DRIVEN • INDEPENDENT • NEUTRAL

finos.org

FINOS hosts over 60 projects across the financial services stack

See all our projects:



SIGs (Communities of practice)



Open Source
READINESS



AI READINESS



DevOps
Automation



DIVERSITY
EQUITY +
INCLUSION
Special Interest Group



Regulation Innovation
SPECIAL INTEREST GROUP

Problem Space

OPEN STANDARDS



OPEN SOURCE SOFTWARE



morphir



Solution Space



The FINOS AI Governance Framework

Open Source AI took the world by storm

WORLD ECONOMIC FORUM

EMERGING TECHNOLOGIES

What is open-source AI and how could DeepSeek change the industry?

Feb 5, 2025

InfoWorld

Topics Spotlight: Prioritizing your AI investments Videos

EU supports AI challenge to Silicon Valley and China

News Feb 03, 2025 • 5 mins

Artificial Intelligence Generative AI

POLITICO

German election War in Ukraine Munich Security Conference Newsletters Podcasts Poll of Polls Policy

NEWS > TECHNOLOGY

Europe's AI hopes rebound after DeepSeek success

New Chinese model suggests winning the AI race could be done on the cheap.

SEARCH

FORTUNE

SIGN IN Subscribe Now

Home News Tech Finance Leadership Well Education Fortune 500

AI

TECH: A.I.

Sam Altman says OpenAI is 'on the wrong side of history' and needs a new open-source strategy after DeepSeek shock

BY AFP February 1, 2025 at 3:51PM GMT

OpenAI CEO Sam Altman during a White House news conference on Jan. 21. ANDREW HARNIK—GETTY IMAGES

CLOUD

Google Cloud donates A2A to Linux Foundation

JUNE 23, 2025

Rao Surapaneni

VP and GM

Business Application Platform

Todd Segal

Principal Engineer

Business Application Platform

Michael Vakoc

Product Manager

Google Cloud

Share

 A2A protocol

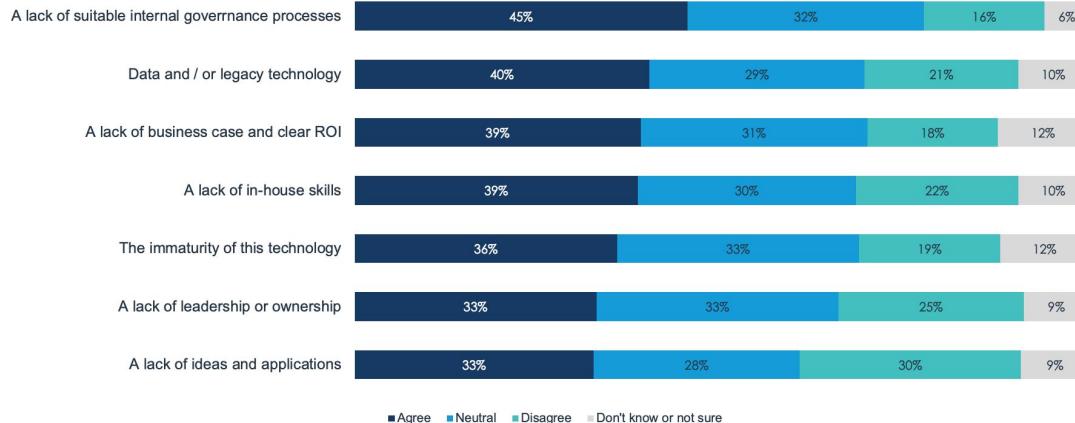


HOT OFF THE PRESS

ONE WORD: GOVERNANCE!

Figure 12

My organization's use of GenAI is limited by, or we do not use GenAI due to: (select one response per row)



2024 FINOS State of Open Source Survey, Q16, Sample Size = 249





An industry wide, broad consensus, truly open framework to govern AI

Why it matters for the industry?

AI is a transformational technology, but the path to adoption is bumpy

Why set up the AI Readiness SIG?

We found many members struggling with the same issues, quite simply getting access to this technology. Many were creating similar solutions to the problem, but efforts were fragmented.

Why open source?

Beyond the 'standard' benefits of open source (pooled resources, rapid development, ...), this is an opportunity to bring FSIs, regulators, technology platform providers into a single place

"A GenAI Governance Framework for FSIs by FSIs"

- *Johnna Powell, Managing Director, and Head of Technology, Research and Innovation at DTCC*

SEC AI Roundtable, 27-03-2025





FINOS AI Governance Framework

A comprehensive collection of risks and mitigations that support on-boarding, development of, and running Generative AI solutions

version: v1 (June 09, 2024)

AI, especially Generative AI, is reshaping financial services, enhancing products, client interactions, and productivity. However, challenges like hallucinations and model unpredictability make safe deployment complex. Rapid advancements require flexible governance.

Financial institutions are eager to adopt AI but face regulatory hurdles. Existing frameworks may not address AI's unique risks, necessitating an adaptive governance model for safe and compliant integration.

The following framework has been developed by [FINOS \(Fintech Open Source Foundation\)](#) members, providing comprehensive catalogue of risks and associated mitigation. We suggest using our [heuristic risk identification framework](#) to determine which risks are most relevant for a given use case.

Risk Catalogue

Operational

[AIR-OP-004](#)

Hallucination and Inaccurate Outputs

SummaryLLM hallucinations occur when a model generates confident but incorrect ...

[Read more](#)

[AIR-OP-005](#)

Foundation Model Versioning

SummaryFoundation model instability refers to unpredictable changes in model behavior ...

[Read more](#)

[AIR-OP-006](#)

Non-Deterministic Behaviour

SummaryLLMs exhibit non-deterministic behaviour, meaning they can generate different outputs ...

[Read more](#)

[AIR-OP-007](#)

Availability of Foundational Model

SummaryFoundation models often rely on GPU-heavy infrastructure hosted by third-party ...

[Read more](#)



AI GOVERNANCE
FRAMEWORK

FINOS incubating

The AI Governance Framework v.1.0

<https://air-governance-framework.finos.org/>



CHECK IT OUT



LSEG



Morgan Stanley



NatWest

publicis
sapient



finos.org



AI GOVERNANCE
FRAMEWORK

FINOS incubating

Stakeholders

For Technical Teams

- Clear guidelines & best practices for GenAI integration
- Early risk identification during development
- Promotes secure, robust GenAI systems

For CISOs & Risk Management

- Provides a 'toolbox' to assess AI security risks early
- Ensures AI aligns with internal policies & regulations

For Heads of AI, CTOs/CIOs

- Aligns AI initiatives with strategy & tech roadmap
- Improves resource optimization & AI ROI

For Vendors & Service Providers

- Clarifies governance expectations of financial institutions
- Builds client trust through adherence to established frameworks



AI GOVERNANCE
FRAMEWORK

FINOS incubating

Structure

- **Risk Catalogue** – groups 17 AI risks under three domains:
 - Operational (OP)**,
 - Security (SEC)** and
 - Regulatory & Compliance (RC)**.

Risk Catalogue					
Operational		Foundation		Non-Deterministic	
AIR-OP-004	Hallucination and Inaccurate Outputs	AIR-OP-005	Foundation Model Versioning	AIR-OP-006	Non-Deterministic Behaviour
Summary:LM hallucinations occur when a model generates confident but incorrect ...	Summary:Foundation model instability refers to unpredictable changes in model behavior ...	Summary:LM exhibits non-deterministic behaviour, meaning it can generate different outputs ...			
Read more	Read more	Read more			
AIR-OP-014	Inadequate System Alignment	AIR-OP-016	Bias and Discrimination	AIR-OP-017	Lack of Explainability
Summary:LM-powered RAG systems may generate responses that diverge from their ...	Summary:AI systems can systematically discriminate through biased training ...	Summary:Systems, particularly those using complex foundation models, often lack ...			
Read more	Read more	Read more			
AIR-OP-019	Data Quality and Drift	AIR-OP-020	Reputational Risk	AIR-OP-018	Model Overreach / Expanded Use
Summary:Generative AI systems rely heavily on the quality and freshness of their ...	Summary:failures or misperceptions—especially in customer-facing systems—can quickly escalate into ...	Summary:Model overreach occurs when AI systems are used beyond their ...			
Read more	Read more	Read more			

AIR-<DOMAIN>-### e.g. **AIR-OP-005**.

- **Mitigation Catalogue** – lists 17 controls, organised as
 - Preventative (PREV)** and
 - Detective (DET)** safeguards.

Mitigation Catalogue					
Preventative			Detective		
AIR-PREV-002	Data Filtering From External Knowledge Bases	AIR-PREV-003	User/App/Model Firewall/Filtering	AIR-PREV-005	System Acceptance Testing
Purposive control addresses the critical need to sanitize, filter, and ...	Protective security for AI systems involves monitoring and filtering interactions ...	Purpose:System Acceptance Testing (SAT) for AI systems is a crucial ...			
Read more	Read more	Read more			
AIR-PREV-007	Legal and Contractual Frameworks for AI Systems	AIR-PREV-008	Quality of Service (QoS) and DDoS Prevention for AI Systems	AIR-PREV-010	Data Quality & Classification/Sensitivity
Purposive legal and contractual agreements are essential for governing the ...	Purposive:the increasing integration of Artificial Intelligence (AI) into financial applications, ...	Purpose:Model Version Pinning is the deliberate practice of selecting and ...			
Read more	Read more	Read more			
AIR-PREV-014	Encryption of AI Data at Rest	AIR-PREV-017	AI Firewall Implementation and Management	AIR-PREV-012	Role-Based Access Control for AI Data
Purposive:encryption of data at rest is a fundamental security control ...	Purposive:AI Firewall is conceptualized as a specialized security system ...	Purposive:Role-Based Access Control (RBAC) is a fundamental security mechanism designed ...			
Read more	Read more	Read more			

finos.org



AI GOVERNANCE
FRAMEWORK

FINOS incubating



Traceability

Risks list the Controls that mitigate them, and Controls list the Risks they cover, plus all linked external clauses—**making it easy to prove “this EU-AI-Act article is handled by these three controls,”** or “this control satisfies these NIST & ISO requirements.” Examples:

- **EU AI Act** – e.g., *III.S2.A9 Risk-Management System, III.S2.A15 Accuracy/Robustness*
- **NIST AI 600-1** – e.g., *2.2 Confabulation, 2.8 Information Integrity*
- **NIST SP 800-53 r5** – e.g., *AC-2 Account Management, AU-2 Event Logging, SI-4 System Monitoring*
- **ISO / IEC 42001** – e.g., *A.6.2.8 Event Logs, A.3.2 Roles & Responsibilities*
- **OWASP LLM Top 10 & ML Top 10, FFIEC IT Handbook** references appear where relevant

OWASP LLM Top 10 References
LLM09:2025 Misinformation

OWASP ML Top 10 References
ML09:2023 Output Integrity Attack

FFIEC References
DAM: III Risk Management of Development, Acquisition, and Maintenance
AUD: Risk Assessment and Risk-Based Auditing
MGT: II Risk Management

EU AI Act References
III.S2.A15: Accuracy, Robustness and Cybersecurity
III.S2.A13: Transparency and Provision of Information to Deployers
III.S2.A9: Risk Management System

NIST AI 600-1 References
2.2. Confabulation

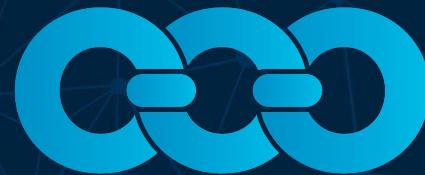
How other FINOS projects benefit from AIGF



BRINGING OUR TECH RISK EFFORTS TOGETHER



**AI GOVERNANCE
FRAMEWORK**



**FINOS
COMMON
CLOUD
CONTROLS**

**The FINOS AI Governance Framework Rule Sets
Combined With Codified Controls**



AI GOVERNANCE
FRAMEWORK



FINOS
COMMON
CLOUD
CONTROLS

= CC4AI

Common Controls for AI Services

What problem are we trying to solve?

Fragmented, proprietary controls are not enough for regulated markets.

Why it matters for you?

Without shared standards, financial institutions face heightened regulatory, security, and concentration risks. Vendors have to comply with a fragmented, highly opinionated variety of per-customer controls.

Why open source?

The FINOS AI Governance Framework demands a lot from cyber-security professionals. Open Source, machine readable, certifiable controls. What if open source could reduce the pain?

"Having played a key role in establishing the AI Readiness programme, we are extremely proud to see FINOS take this major step forward towards establishing a shared approach to AI governance. By collaborating across our industry on a common controls framework we make everyone stronger"

- Ian Micallef, MD,
Developer Enablement, Citi



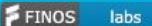


How To Get Involved

AI Readiness SIG

A space for financial institutions to address GenAI-related adoption challenges. AIGF was incubated here and Responsibile AI WG is next.

Applied GenAI Group



Exploring the potential of GenAI in FSIs and foster the development of cross-industry experimentation focused on pre-competitive use cases

AI Fund Governing Board



AI decision makers body to validate industry needs and direct pooled funding towards Academia & SMEs to accelerate FINOS open source AI projects



GET IN TOUCH



PARTICIPATE TO APPLIED GENAI



ADOPT THE FRAMEWORK



Thank You!

